

In [3]:

```
!wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/84.0.4147.105 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-GB,en-US;q=0.9,en;q=0.8" --header="Referer: https://www.kaggle.com/" "https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/11835/224935/compressed/train.csv.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1597640517&Signature=QRdQpSVbTEc2jGyNRotvi7hYO%2BEuWfNLBjmfVggVHIOPyKxBu8QENQovyEOEeUCF7C7Nuu2FjpI2jKhLYCuh18NSRliAC4iGcTWKYG9oBkqA3G0cnjm2vQWnI820mboarvJoyU8rU5%2BTwm4qTdLq9g3a77J3yH%2FwOstZ53cjhnzPTM208EjyLzFPZTOqzqieEilWZB9GKBbFXIjsn6CAFpod437PNHiPRyoLbVOi3kJ8BruAI6gDu0pXf8jCaCSQCJGml0oYPGqVss15NkrwlJcGUHvyVhfKt7S7V%2Fcp%2Br0MZ7p1puJW2F7IYO%2BuQi%2BMJvHxOclbnKRJmaszqiM7gw%3D%3D&response-content-disposition=attachment%3B+filename%3Dtrain.csv.zip" -c -O 'train.csv.zip'
```

```
--2020-08-14 23:24:07-- https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/11835/224935/compressed/train.csv.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1597640517&Signature=QRdQpSVbTEc2jGyNRotvi7hYO%2BEuWfNLBjmfVggVHIOPyKxBu8QENQovyEOEeUCF7C7Nuu2FjpI2jKhLYCuh18NSRliAC4iGcTWKYG9oBkqA3G0cnjm2vQWnI820mboarvJoyU8rU5%2BTwm4qTdLq9g3a77J3yH%2FwOstZ53cjhnzPTM208EjyLzFPZTOqzqieEilWZB9GKBbFXIjsn6CAFpod437PNHiPRyoLbVOi3kJ8BruAI6gDu0pXf8jCaCSQCJGml0oYPGqVss15NkrwlJcGUHvyVhfKt7S7V%2Fcp%2Br0MZ7p1puJW2F7IYO%2BuQi%2BMJvHxOclbnKRJmaszqiM7gw%3D%3D&response-content-disposition=attachment%3B+filename%3Dtrain.csv.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.142.128, 74.125.195.128, 74.125.20.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.142.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 456337398 (435M) [application/zip]
Saving to: 'train.csv.zip'
```

```
train.csv.zip      100%[=====>] 435.20M  61.9MB/s    in 7.0s
```

```
2020-08-14 23:24:14 (61.8 MB/s) - 'train.csv.zip' saved [456337398/456337398]
```

In [4]:

```
!unzip train.csv.zip
```

```
Archive:  train.csv.zip
  inflating: train.csv
```

In [31]:

```
# Importing required libraries.
import warnings
import itertools
import numpy as np
import pandas as pd
import seaborn as sns
import lightgbm as lgb
import matplotlib.pyplot as plt
from tqdm import tqdm_notebook as tqdm
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, log_loss
import dask.dataframe as dd
import dask
import gc
warnings.filterwarnings(action='ignore')
sns.set_style('whitegrid')
```

In [32]:

```
# This is to be used for memory optimization because the data is very large.
# For more information about iinfo refer- https://numpy.org/doc/stable/reference/generate-d/numpy.iinfo.html
```

```

# np.iinfo gives machine limits for integer types, so based on max and minimum value of that feature we can \
# convert the feature values into best suitable optimum datatype to reduce memory usage.
def reduce_mem_usage(df):

```

```

    start_mem = df.memory_usage().sum() / 1024**2
    print('Memory usage of dataframe is {:.2f} MB'.format(start_mem))

    for col in df.columns:
        col_type = df[col].dtype

        if col_type != object:
            c_min = df[col].min()
            c_max = df[col].max()
            if str(col_type)[:3] == 'int':
                if c_min > np.iinfo(np.int8).min and c_max < np.iinfo(np.int8).max:
                    df[col] = df[col].astype(np.int8)
                elif c_min > np.iinfo(np.int16).min and c_max < np.iinfo(np.int16).max:
                    df[col] = df[col].astype(np.int16)
                elif c_min > np.iinfo(np.int32).min and c_max < np.iinfo(np.int32).max:
                    df[col] = df[col].astype(np.int32)
                elif c_min > np.iinfo(np.int64).min and c_max < np.iinfo(np.int64).max:
                    df[col] = df[col].astype(np.int64)
            else:
                if c_min > np.finfo(np.float16).min and c_max < np.finfo(np.float16).max:
                    df[col] = df[col].astype(np.float16)
                elif c_min > np.finfo(np.float32).min and c_max < np.finfo(np.float32).max:
                    df[col] = df[col].astype(np.float32)
                else:
                    df[col] = df[col].astype(np.float64)
        else:
            df[col] = df[col].astype('category')

    end_mem = df.memory_usage().sum() / 1024**2
    print('Memory usage after optimization is: {:.2f} MB'.format(end_mem))
    print('Decreased by {:.1f}%'.format(100 * (start_mem - end_mem) / start_mem))

    return df

```

```

# In featureModify I encoded the output event of a pilot's state \
# while going through certain experiment.

```

```

def featureModify(isTrain, numRows):

    if isTrain:
        df = dd.read_csv('train.csv',nrows=numRows)
        df = df.compute()
        df = reduce_mem_usage(df)
        df['event'] = df['event'].map({
            'A':0,
            'B':1,
            'C':2,
            'D':3
        })
    else:
        df = dd.read_csv('test.csv',nrows=numRows)
        df = df.compute()
        df = reduce_mem_usage(df)

    return df

```

In [35]:

```

# Here we can see the 1.076 GB memory usage was reduced to 278.52 MB so its good for better usage.
train = featureModify(True, None)
y = train['event']

```

```
train = train.drop('event',axis=1)
print(train.shape)
print(train.columns)
```

Memory usage of dataframe is 1076.93 MB

Memory usage after optimization is: 278.52 MB

Decreased by 74.1%

(4867421, 27)

```
Index(['crew', 'experiment', 'time', 'seat', 'eeg_fp1', 'eeg_f7', 'eeg_f8',
      'eeg_t4', 'eeg_t6', 'eeg_t5', 'eeg_t3', 'eeg_fp2', 'eeg_o1', 'eeg_p3',
      'eeg_pz', 'eeg_f3', 'eeg_fz', 'eeg_f4', 'eeg_c4', 'eeg_p4', 'eeg_poz',
      'eeg_c3', 'eeg_cz', 'eeg_o2', 'ecg', 'r', 'gsr'],
      dtype='object')
```

In []:

```
# Test data's memory usage was reduced from 3.974 GB to 1.079 GB
test = featureModify(False,None)
print(test.shape)
print(test.columns)
```

Memory usage of dataframe is 3974.83 MB

Memory usage after optimization is: 1079.37 MB

Decreased by 72.8%

(17965143, 28)

```
Index(['id', 'crew', 'experiment', 'time', 'seat', 'eeg_fp1', 'eeg_f7',
      'eeg_f8', 'eeg_t4', 'eeg_t6', 'eeg_t5', 'eeg_t3', 'eeg_fp2', 'eeg_o1',
      'eeg_p3', 'eeg_pz', 'eeg_f3', 'eeg_fz', 'eeg_f4', 'eeg_c4', 'eeg_p4',
      'eeg_poz', 'eeg_c3', 'eeg_cz', 'eeg_o2', 'ecg', 'r', 'gsr'],
      dtype='object')
```

Training on the model with best results: LightGBM

In []:

```
def pipeline(train,test):

    print('Splitting the data into train and cross_validation')
    X_train, X_cv, y_train, y_cv = train_test_split(train, y, test_size=0.25, shuffle=False)

    train = lgb.Dataset(X_train, label = y_train, categorical_feature = [1])
    print('Done')

    print('Making the train and cv data from LightGBM model')
    del y_train
    # gc is garbage collector it returns how many objects it has collected and deallocated.
    gc.collect()
    cv = lgb.Dataset(X_cv, label = y_cv, categorical_feature = [1])
    del y_cv
    gc.collect()
    print('Done')

    # Declaring the hyperparameter values.
    params = {
        "objective" : "multiclass", # used for multiclass softmax classifier
        "metric" : "multi_error", # Error rate for multiclass classification
        "boosting" : 'gbdt', # Using Gradient Boosted Decision Trees
        "num_class":4, # Number of desired output classes is 4
        "num_leaves" : 30, # Number of leaves in Tree based algorithms
        "learning_rate" : 0.01,
        "bagging_fraction" : 0.9, # This is randomly select 90% of data without resampling

        "bagging_seed" : 0, # it will decrease impact of high variance on data
        "num_threads" : 4, # Random seeds for bagging
```

```

        "colsample_bytree" : 0.5,      # Subsampling fraction for feature
        'min_data_in_leaf':100,       # Threshold on Data in a leaf
        'min_split_gain':0.00019     # Minimum gain threshold for splitting the node
    }

# Training Stage:
print('Training Started:')
model = lgb.train(
    params,
    train_set = train,
    num_boost_round=2000,
    early_stopping_rounds=200,
    verbose_eval=100,
    valid_sets=[train,cv]
)

print('Training Done!')

# Exporting the trained model.
import joblib
joblib.dump(model, 'final_model.pkl')

print('Predicting the output probabilities for each event for test data:')
# Saving the ids for further merging of features in csv.
df_sub = pd.DataFrame()
df_sub['id'] = test['id']
test = test.drop('id',axis=1)

# Predicting the probability values for each event.
y_pred = model.predict(test,num_iteration=model.best_iteration)
df_sub = pd.DataFrame(np.concatenate((np.arange(len(test))[:, np.newaxis], y_pred), axis=1), columns=['id', 'A', 'B', 'C', 'D'])
print('Predictions Done')

# Creating the submission csv file.
df_sub['id'] = df_sub['id'].astype(int)
print(df_sub)
df_sub.to_csv("predictions.csv", index=False)

return y_pred

```

```
predictions = pipeline(train,test)
```

Splitting the data into train and cross_validation

Done

Making the train and cv data from LightGBM model

Done

Training Started:

Training until validation scores don't improve for 200 rounds.

[100] training's multi_error: 0.0732298 valid_1's multi_error: 0.150007

[200] training's multi_error: 0.0513315 valid_1's multi_error: 0.0836788

[300] training's multi_error: 0.0446361 valid_1's multi_error: 0.0819834

[400] training's multi_error: 0.0376761 valid_1's multi_error: 0.0810803

[500] training's multi_error: 0.0339035 valid_1's multi_error: 0.0818256

Early stopping, best iteration is:

[392] training's multi_error: 0.0380341 valid_1's multi_error: 0.08098

Training Done!

Predicting the output probabilities for each event for test data:

Predictions Done

	id	A	B	C	D
0	0	0.973573	0.002153	0.020307	0.003967
1	1	0.969859	0.002878	0.023126	0.004137

2	0.973605	0.002191	0.020116	0.004088
3	0.969847	0.002912	0.023105	0.004136
4	0.974038	0.002179	0.019839	0.003944
...
17965138	17965138	0.882729	0.003132	0.106111
17965139	17965139	0.941786	0.002305	0.048648
17965140	17965140	0.882740	0.003190	0.106061
17965141	17965141	0.941543	0.002313	0.048936
17965142	17965142	0.883312	0.003137	0.105549

[17965143 rows x 5 columns]

This is the kaggle score which I got.I got the public score of 0.31120 and orivate score of 0.55509

Model_building5_LGBop.zip

5 days ago by AtharvaMusale

[add submission details](#)

0.55509

0.31120



Predictions-

In [4]:

```
import joblib
model = joblib.load('/content/drive/My Drive/ML case study/models/final_model_cs.pkl')
```

In [22]:

```
def function1(X):
    pred = model.predict(X,num_iteration=model.best_iteration)
    return pred
```

In [36]:

```
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score,log_loss,confusion_matrix
lb = LabelEncoder()
def function2(train,y):
    print('Testing the model for Set of Datapoints:')
    print('#'*40)
    X_train,X_cv,y_train,y_cv = train_test_split(train,y,test_size=0.2,random_state=30)
    y_pred = function1(X_cv)
    y_pred = np.argmax(y_pred, axis=1)
    print('Accuracy score of a model is:',accuracy_score(y_cv,y_pred))
    print('Confusion Matrix of a predictions is:\n',confusion_matrix(y_cv,y_pred))
    print('*'*100)

    print('Testing the model for single input:')
    print('#'*35)
    train['experiment'] = lb.fit_transform(train['experiment'])
    index = np.random.randint(0,30000)
    data = train.iloc[index,:]
    print('Data to be used for prediction is:\n',data)
    print('#'*35)
    y_pred = function1(np.array(data).reshape(1,-1))
    print('Predicted output event is:',np.argmax(y_pred, axis=1))
    ytrue =(y.iloc[index])
    print('Actual output event is:',ytrue)

    print('*'*100)
```

function2(train,y)

```
Testing the model for Set of Datapoints:
#####
Accuracy score of a model is: 0.9512565678978104
Confusion Matrix of a predictions is:
[[567721  181    375   10151
```

```
[ 307721 101 373 1013]
[ 9692 16371 0 0]
[ 72 0 330813 0]
[ 36116 0 0 11129]]
*****
*****
Testing the model for single input:
#####
Data to be used for prediction is:
crew 1.000000
experiment 0.000000
time 10.328125
seat 0.000000
eeg_fp1 4.082031
eeg_f7 -5.148438
eeg_f8 -2.011719
eeg_t4 -0.316162
eeg_t6 2.355469
eeg_t5 -9.617188
eeg_t3 6.359375
eeg_fp2 11.187500
eeg_o1 -1.160156
eeg_p3 -9.742188
eeg_pz -7.375000
eeg_f3 -8.859375
eeg_fz -3.935547
eeg_f4 -2.685547
eeg_c4 -3.185547
eeg_p4 -4.679688
eeg_poz -3.886719
eeg_c3 -15.039062
eeg_cz -9.382812
eeg_o2 1.080078
ecg -18352.000000
r 666.000000
gsr 835.000000
Name: 1144, dtype: float64
#####
Predicted output event is: [2]
Actual output event is: 2
*****
*****
```

In []: