

So here is what I am going to try in this notebook , I am going to drop experiment feature and keep the time feature as it is. I am simply going to train my model on this data as it is and see the results on the best model.

In []:

```
!wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/84.0.4147.105 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-GB,en-US;q=0.9,en;q=0.8" --header="Referer: https://www.kaggle.com/" "https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/11835/224935/compressed/train.csv.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1597053235&Signature=LPrAGyUvFAFA7ZkkwUWyc0srI5488quzzTLdj1ThHt3H2isET80a8V6dHQQRuHzZHJgRq%2FBF%2BGzLxf0DEaB0HB0iYF2AnxlKcxEU3xbWcdZF9at7J8HOfFNB9t3AezeghzCnu6FAcLvKRTHVJrCJnLBQk6wQgFHXgJPWh3dQeuhnJwEVpkj2rFqBkN2Ray3r2WZeL01xO6VZN807zGWg4ORaQikb3I12Vid%2BrN5ubxeljOQ3UfDJ%2B8hDMVMSVu86SHZY6xREyLN5dvMn%2BLK8j98Gzw4a5dXx3F3WdG13t5a1%2F%2BPB2H54Ge4j6%2F15AB0gwf%2FWmcgi7%2ByDeNB9xAiMDw%3D%3D&response-content-disposition=attachment%3B+filename%3Dtrain.csv.zip" -c -O 'train.csv.zip'
```

```
--2020-08-08 11:48:18-- https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/11835/224935/compressed/train.csv.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1597053235&Signature=LPrAGyUvFAFA7ZkkwUWyc0srI5488quzzTLdj1ThHt3H2isET80a8V6dHQQuHzZHJgRq%2FbF%2BGzLxf0DEaB0HB0iYF2AnxlKcxEU3xbWcdZF9at7J8HOfFNB9t3AezeghzCnu6FAcLvKRTHVJRcJnLBQk6wQgFHxgJPwH3dQeuhnJwEVpkj2rFqBkN2Ray3r2WZeL01xO6VZN807zGWg4ORaQikb3I12Vid%2BrN5ubxeljOQ3UfDJ%2B8hDMVMSVu86SHZY6xREyLN5dvMn%2BLK8j98Gzw4a5dXx3F3WdGl3t5a1%2F%2BBP2H54Ge4j6%2F15AB0gwf%2FWmcgi7%2ByDeNB9xAiMDw%3D%3D&response-content-disposition=attachment%3B+filename%3Dtrain.csv.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 173.194.216.128, 173.194.217.128, 172.217.193.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|173.194.216.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 456337398 (435M) [application/zip]
Saving to: 'train.csv.zip'
```

```
train.csv.zip      100%[=====>] 435.20M   119MB/s   in 3.6s
```

```
2020-08-08 11:48:22 (119 MB/s) - 'train.csv.zip' saved [456337398/456337398]
```

In []:

```
! unzip train.csv.zip
```

```
Archive:  train.csv.zip
  inflating: train.csv
```

In []:

```
import warnings
import itertools
import numpy as np
import pandas as pd
import seaborn as sns
import lightgbm as lgb
import matplotlib.pyplot as plt
from tqdm import tqdm_notebook as tqdm
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, log_loss
import dask
import gc
```

```
from yellowbrick.text import TSNEVisualizer
```

```
%matplotlib inline
plt.style.use("fivethirtyeight")
```

```
warnings.filterwarnings(action='ignore')
sns.set_style('whitegrid')
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
```

```
import pandas.util.testing as tm
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:144: FutureWarning: The sklearn.metrics.classification module is deprecated in version 0.22 and will be removed in version 0.24. The corresponding classes / functions should instead be imported from sklearn.metrics. Anything that cannot be imported from sklearn.metrics is now part of the private API.
warnings.warn(message, FutureWarning)
```

```
In [ ]:
```

```
# This is to be used for memory optimization because the data is very large.
def reduce_mem_usage(df):
    """ iterate through all the columns of a dataframe and modify the data type
    to reduce memory usage.
    """
    start_mem = df.memory_usage().sum() / 1024**2
    print('Memory usage of dataframe is {:.2f} MB'.format(start_mem))

    for col in df.columns:
        col_type = df[col].dtype

        if col_type != object:
            c_min = df[col].min()
            c_max = df[col].max()
            if str(col_type)[:3] == 'int':
                if c_min > np.iinfo(np.int8).min and c_max < np.iinfo(np.int8).max:
                    df[col] = df[col].astype(np.int8)
                elif c_min > np.iinfo(np.int16).min and c_max < np.iinfo(np.int16).max:
                    df[col] = df[col].astype(np.int16)
                elif c_min > np.iinfo(np.int32).min and c_max < np.iinfo(np.int32).max:
                    df[col] = df[col].astype(np.int32)
                elif c_min > np.iinfo(np.int64).min and c_max < np.iinfo(np.int64).max:
                    df[col] = df[col].astype(np.int64)
            else:
                if c_min > np.finfo(np.float16).min and c_max < np.finfo(np.float16).max:
                    df[col] = df[col].astype(np.float16)
                elif c_min > np.finfo(np.float32).min and c_max < np.finfo(np.float32).max:
                    df[col] = df[col].astype(np.float32)
                else:
                    df[col] = df[col].astype(np.float64)
        else:
            df[col] = df[col].astype('category')

    end_mem = df.memory_usage().sum() / 1024**2
    print('Memory usage after optimization is: {:.2f} MB'.format(end_mem))
    print('Decreased by {:.1f}%'.format(100 * (start_mem - end_mem) / start_mem))

    return df

def featureModify(isTrain, numRows):
    if isTrain:
        df = dd.read_csv('train.csv',nrows=numRows)
        df = df.compute()
        # df['pilot'] = 100*df['crew']+df['seat']
        df = reduce_mem_usage(df)
        df['event'] = df['event'].map({
            'A':0,
            'B':1,
            'C':2,
            'D':3
        })
    else:
        df = dd.read_csv('test.csv',nrows=numRows)
        df = df.compute()
        # df['pilot'] = 100*df['crew']+df['seat']
```

```
df = reduce_mem_usage(df)
```

```
    return df
# train = featureModify(True, None)
# y = train['event']
# train = train.drop('event',axis=1)
# print(train.shape)
# print(train.columns)
```

In []:

```
# I didnt do this for LightGBM
train = train.drop(columns=['experiment'],axis=1)
```

In []:

```
col = train.columns
```

In []:

```
# I have used this normalisation only for naive bayes
# from sklearn.preprocessing import MinMaxScaler
# ms = MinMaxScaler()
# train = ms.fit_transform(train)
```

In []:

```
train = pd.DataFrame(train)
```

In []:

```
train.columns = col
```

In []:

```
train.head()
```

Out[]:

	crew	time	seat	eeg_fp1	eeg_f7	eeg_f8	eeg_t4	eeg_t6	eeg_t5	eeg_t3	eeg_fp2	eeg_o1	
0	1	0.011719	1	-5.285156	26.781250	-9.523438	12.796875	16.718750	33.75000	23.718750	6.695312	29.234375	24
1	1	0.015625	1	-2.427734	28.437500	-9.320312	-3.757812	15.968750	30.43750	21.015625	6.476562	26.640625	24
2	1	0.019531	1	10.671875	30.421875	15.351562	24.718750	16.140625	32.15625	25.437500	0.088684	28.125000	26
3	1	0.023438	1	11.453125	25.609375	2.433594	12.414062	20.531250	31.50000	19.140625	0.256592	30.656250	24
4	1	0.027344	1	7.285156	25.937500	0.113586	5.746094	19.828125	28.75000	20.578125	1.953125	31.718750	25

Naive Bayes

In []:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(train,y,test_size =0.2, random_state=40
)
```

In []:

```
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
```

```
print(y_test.shape)
```

```
(3893936, 26)
(973485, 26)
(3893936,)
(973485,)
```

```
In [ ]:
```

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import RandomizedSearchCV
param_dist = {"alpha": [0.0000001, 0.000001, 0.00001, 0.0001, 0.001, 0.01, 0.1, 1]}
nb = MultinomialNB()
rsc = RandomizedSearchCV(nb, param_dist, cv=4, verbose=500)
rsc.fit(X_train, y_train)
print("Tuned Multinomial Naive Bayes Parameters: {}".format(rsc.best_params_))
print("Best score is {}".format(rsc.best_score_))
```

```
Fitting 4 folds for each of 8 candidates, totalling 32 fits
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[CV] alpha=1e-07 ..... alpha=1e-07, score=0.585, total= 2.5s
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 2.5s remaining: 0.0s
[CV] alpha=1e-07 ..... alpha=1e-07, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 2 out of 2 | elapsed: 4.9s remaining: 0.0s
[CV] alpha=1e-07 ..... alpha=1e-07, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 3 out of 3 | elapsed: 7.3s remaining: 0.0s
[CV] alpha=1e-07 ..... alpha=1e-07, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 4 out of 4 | elapsed: 9.7s remaining: 0.0s
[CV] alpha=1e-06 ..... alpha=1e-06, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 5 out of 5 | elapsed: 12.2s remaining: 0.0s
[CV] alpha=1e-06 ..... alpha=1e-06, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 6 out of 6 | elapsed: 14.6s remaining: 0.0s
[CV] alpha=1e-06 ..... alpha=1e-06, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 7 out of 7 | elapsed: 17.0s remaining: 0.0s
[CV] alpha=1e-06 ..... alpha=1e-06, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 8 out of 8 | elapsed: 19.5s remaining: 0.0s
[CV] alpha=1e-05 ..... alpha=1e-05, score=0.585, total= 2.5s
[Parallel(n_jobs=1)]: Done 9 out of 9 | elapsed: 21.9s remaining: 0.0s
[CV] alpha=1e-05 ..... alpha=1e-05, score=0.585, total= 2.5s
[Parallel(n_jobs=1)]: Done 10 out of 10 | elapsed: 24.4s remaining: 0.0s
[CV] alpha=1e-05 ..... alpha=1e-05, score=0.585, total= 2.5s
[Parallel(n_jobs=1)]: Done 11 out of 11 | elapsed: 26.8s remaining: 0.0s
[CV] alpha=1e-05 ..... alpha=1e-05, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 12 out of 12 | elapsed: 29.3s remaining: 0.0s
[CV] alpha=0.0001 ..... alpha=0.0001, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 13 out of 13 | elapsed: 31.7s remaining: 0.0s
[CV] alpha=0.0001 ..... alpha=0.0001, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 14 out of 14 | elapsed: 34.1s remaining: 0.0s
[CV] alpha=0.0001 ..... alpha=0.0001, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 15 out of 15 | elapsed: 36.6s remaining: 0.0s
[CV] alpha=0.0001 ..... alpha=0.0001, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 16 out of 16 | elapsed: 39.0s remaining: 0.0s
[CV] alpha=0.001 ..... alpha=0.001, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 17 out of 17 | elapsed: 41.4s remaining: 0.0s
[CV] alpha=0.001 ..... alpha=0.001, score=0.585, total= 2.4s
```

```

[CV] ..... alpha=0.001, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 18 out of 18 | elapsed: 43.8s remaining: 0.0s
[CV] alpha=0.001 .....
[CV] ..... alpha=0.001, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 19 out of 19 | elapsed: 46.2s remaining: 0.0s
[CV] alpha=0.001 .....
[CV] ..... alpha=0.001, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 20 out of 20 | elapsed: 48.6s remaining: 0.0s
[CV] alpha=0.01 .....
[CV] ..... alpha=0.01, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 21 out of 21 | elapsed: 51.1s remaining: 0.0s
[CV] alpha=0.01 .....
[CV] ..... alpha=0.01, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 22 out of 22 | elapsed: 53.5s remaining: 0.0s
[CV] alpha=0.01 .....
[CV] ..... alpha=0.01, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 23 out of 23 | elapsed: 55.9s remaining: 0.0s
[CV] alpha=0.01 .....
[CV] ..... alpha=0.01, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 24 out of 24 | elapsed: 58.3s remaining: 0.0s
[CV] alpha=0.1 .....
[CV] ..... alpha=0.1, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 25 out of 25 | elapsed: 1.0min remaining: 0.0s
[CV] alpha=0.1 .....
[CV] ..... alpha=0.1, score=0.585, total= 2.5s
[Parallel(n_jobs=1)]: Done 26 out of 26 | elapsed: 1.1min remaining: 0.0s
[CV] alpha=0.1 .....
[CV] ..... alpha=0.1, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 27 out of 27 | elapsed: 1.1min remaining: 0.0s
[CV] alpha=0.1 .....
[CV] ..... alpha=0.1, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 28 out of 28 | elapsed: 1.1min remaining: 0.0s
[CV] alpha=1 .....
[CV] ..... alpha=1, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 29 out of 29 | elapsed: 1.2min remaining: 0.0s
[CV] alpha=1 .....
[CV] ..... alpha=1, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 30 out of 30 | elapsed: 1.2min remaining: 0.0s
[CV] alpha=1 .....
[CV] ..... alpha=1, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 31 out of 31 | elapsed: 1.3min remaining: 0.0s
[CV] alpha=1 .....
[CV] ..... alpha=1, score=0.585, total= 2.4s
[Parallel(n_jobs=1)]: Done 32 out of 32 | elapsed: 1.3min remaining: 0.0s
[Parallel(n_jobs=1)]: Done 32 out of 32 | elapsed: 1.3min finished
Tuned Multinomial Naive Bayes Parameters: {'alpha': 1e-07}
Best score is 0.5852497318908169

```

In []:

```

from sklearn.metrics import accuracy_score, log_loss, confusion_matrix
clf = MultinomialNB(alpha=1e-07)
clf = clf.fit(X_train, y_train)
y_pred = clf.predict_proba(X_test)
print(log_loss(y_test, y_pred))

```

0.9223443374768153

Logistic Regression

In []:

```

from sklearn.linear_model import LogisticRegression
from tqdm.auto import tqdm
from sklearn.model_selection import RandomizedSearchCV
# from sklearn.multiclass import OneVsRestClassifier
param_dist = {"C" : [0.001, 0.01, 0.1, 1, 10]}
lr = LogisticRegression(multi_class='ovr')
rsc = RandomizedSearchCV(lr, param_dist, cv=2, verbose=500, n_jobs=1)
rsc.fit(X_train, y_train)

```

```
print("Tuned Logistic Regression Parameters: {}".format(rsc.best_params_))
print("Best score is {}".format(rsc.best_score_))
```

```
Fitting 2 folds for each of 5 candidates, totalling 10 fits
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[CV] C=0.001 .....
[CV] ..... C=0.001, score=0.583, total= 2.4min
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 2.4min remaining: 0.0s
[CV] C=0.001 .....
[CV] ..... C=0.001, score=0.583, total= 2.5min
[Parallel(n_jobs=1)]: Done 2 out of 2 | elapsed: 4.9min remaining: 0.0s
[CV] C=0.01 .....
[CV] ..... C=0.01, score=0.583, total= 2.4min
[Parallel(n_jobs=1)]: Done 3 out of 3 | elapsed: 7.3min remaining: 0.0s
[CV] C=0.01 .....
[CV] ..... C=0.01, score=0.582, total= 2.4min
[Parallel(n_jobs=1)]: Done 4 out of 4 | elapsed: 9.7min remaining: 0.0s
[CV] C=0.1 .....
[CV] ..... C=0.1, score=0.582, total= 2.4min
[Parallel(n_jobs=1)]: Done 5 out of 5 | elapsed: 12.1min remaining: 0.0s
[CV] C=0.1 .....
[CV] ..... C=0.1, score=0.583, total= 2.4min
[Parallel(n_jobs=1)]: Done 6 out of 6 | elapsed: 14.5min remaining: 0.0s
[CV] C=1 .....
[CV] ..... C=1, score=0.582, total= 2.4min
[Parallel(n_jobs=1)]: Done 7 out of 7 | elapsed: 16.9min remaining: 0.0s
[CV] C=1 .....
[CV] ..... C=1, score=0.583, total= 2.4min
[Parallel(n_jobs=1)]: Done 8 out of 8 | elapsed: 19.3min remaining: 0.0s
[CV] C=10 .....
[CV] ..... C=10, score=0.582, total= 2.4min
[Parallel(n_jobs=1)]: Done 9 out of 9 | elapsed: 21.7min remaining: 0.0s
[CV] C=10 .....
[CV] ..... C=10, score=0.582, total= 2.4min
[Parallel(n_jobs=1)]: Done 10 out of 10 | elapsed: 24.1min remaining: 0.0s
[Parallel(n_jobs=1)]: Done 10 out of 10 | elapsed: 24.1min finished
Tuned Logistic Regression Parameters: {'C': 0.01}
Best score is 0.5826544144536531
```

In []:

```
clf = LogisticRegression(multi_class='ovr',C = 0.01)
clf = clf.fit(X_train,y_train)
pred= clf.predict_proba(X_test)
print(log_loss(y_test,pred))
```

0.9206174913509789

Decision Tree

In []:

```
from sklearn.tree import DecisionTreeClassifier
from tqdm.auto import tqdm
max_depth = [10,100,200,300]
for i in tqdm(max_depth):
    print('depth is',i)
    clf = DecisionTreeClassifier(max_depth=i)
    clf.fit(X_train,y_train)
    y_pred = clf.predict_proba(X_test)
    print(log_loss(y_test,y_pred))
```

```
depth is 10
0.49245480396760244
depth is 100
0.04200774665411109
depth is 200
0.04179486955958013
depth is 300
```

```
depth is 300
0.042185144232886884
```

Okay so we observed that the optimum depth size is 200 since it gives us least loss

I am taking the reference of the hyperparamter tuning form the previous model building notebook and training the data to see which one gives best result

In []:

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import RandomizedSearchCV
param_dist = {"max_depth" : [200],
              "criterion" : ['entropy'],
              "random_state" : [40],
              "max_features" : ['auto'],
              "max_leaf_nodes" : [1500,2000,2500]
              }
dtc = DecisionTreeClassifier()
rsc = RandomizedSearchCV(dtc,param_dist,cv=2,verbose=500)
rsc.fit(X_train,y_train)
print("Tuned Decision Tree Parameters: {}".format(rsc.best_params_))
print("Best score is {}".format(rsc.best_score_))
```

```
Fitting 2 folds for each of 3 candidates, totalling 6 fits
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[CV] random_state=40, max_leaf_nodes=1500, max_features=auto, max_depth=200, criterion=entropy
[CV] random_state=40, max_leaf_nodes=1500, max_features=auto, max_depth=200, criterion=entropy, score=0.888, total= 22.2s
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 22.2s remaining: 0.0s
[CV] random_state=40, max_leaf_nodes=1500, max_features=auto, max_depth=200, criterion=entropy
[CV] random_state=40, max_leaf_nodes=1500, max_features=auto, max_depth=200, criterion=entropy, score=0.873, total= 22.9s
[Parallel(n_jobs=1)]: Done 2 out of 2 | elapsed: 45.0s remaining: 0.0s
[CV] random_state=40, max_leaf_nodes=2000, max_features=auto, max_depth=200, criterion=entropy
[CV] random_state=40, max_leaf_nodes=2000, max_features=auto, max_depth=200, criterion=entropy, score=0.898, total= 23.2s
[Parallel(n_jobs=1)]: Done 3 out of 3 | elapsed: 1.1min remaining: 0.0s
[CV] random_state=40, max_leaf_nodes=2000, max_features=auto, max_depth=200, criterion=entropy
[CV] random_state=40, max_leaf_nodes=2000, max_features=auto, max_depth=200, criterion=entropy, score=0.884, total= 23.4s
[Parallel(n_jobs=1)]: Done 4 out of 4 | elapsed: 1.5min remaining: 0.0s
[CV] random_state=40, max_leaf_nodes=2500, max_features=auto, max_depth=200, criterion=entropy
[CV] random_state=40, max_leaf_nodes=2500, max_features=auto, max_depth=200, criterion=entropy, score=0.905, total= 23.8s
[Parallel(n_jobs=1)]: Done 5 out of 5 | elapsed: 1.9min remaining: 0.0s
[CV] random_state=40, max_leaf_nodes=2500, max_features=auto, max_depth=200, criterion=entropy
[CV] random_state=40, max_leaf_nodes=2500, max_features=auto, max_depth=200, criterion=entropy, score=0.891, total= 24.2s
[Parallel(n_jobs=1)]: Done 6 out of 6 | elapsed: 2.3min remaining: 0.0s
[Parallel(n_jobs=1)]: Done 6 out of 6 | elapsed: 2.3min finished
Tuned Decision Tree Parameters: {'random_state': 40, 'max_leaf_nodes': 2500, 'max_features': 'auto', 'max_depth': 200, 'criterion': 'entropy'}
Best score is 0.8984171799433787
```

In []:

```
clf = DecisionTreeClassifier(max_depth = 200, criterion= 'entropy', max_leaf_nodes = None, random_state = 40)
clf = clf.fit(X_train,y_train)
pred = clf.predict_proba(X_test)
```

NameError

Traceback (most recent call last)

```
<ipython-input-14-8faef9408baf> in <module>()
      2 clf = clf.fit(X_train,y_train)
      3 pred = clf.predict_proba(X_test)
----> 4 print(log_loss(y_cv,pred))
```

NameError: name 'y_cv' is not defined

In []:

```
print(log_loss(y_test,pred))
```

0.021429627516119047

In []:

```
from sklearn.tree import DecisionTreeClassifier
clf = DecisionTreeClassifier(max_depth = 200,criterion = 'entropy', random_state=40)
clf.fit(X_train,y_train)
```

Out[]:

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',
                        max_depth=200, max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort='deprecated',
                        random_state=40, splitter='best')
```

In []:

```
import joblib
# joblib.dump(clf,'decision_tree.pkl')
```

Out[]:

['decision_tree.pkl']

In []:

```
import joblib
clf = joblib.load('decision_tree.pkl')
# clf.predict_pro
```

In []:

```
y_pred = clf.predict_proba(test)
```

LightGBM

For Lightgbm I am not dropping the experiment feature since the I tried it by dropping that feature but I didnt get very good result.

In []:

```
train, train_test, y, y_test = train_test_split(train, y, test_size=0.25, shuffle=False)
train = lgb.Dataset(train, label=y,categorical_feature=[1])
del y
gc.collect()
```

```
train_test = lgb.Dataset(train_test, label=y_test,categorical_feature=[1])
del y_test
gc.collect()
```

Out[]:

0

In []:


```

params = {
    "objective" : "multiclass",
    "metric" : "multi_error",
    'num_class':4,
    "num_leaves" : 30,
    "learning_rate" : 0.01,
    "bagging_fraction" : 0.9,
    "bagging_seed" : 0,
    "num_threads" : 4,
    "colsample_bytree" : 0.5,
    'min_data_in_leaf':100,
    'min_split_gain':0.00019
}
model = lgb.train( params,
                   train_set = train,
                   num_boost_round=2000,
                   early_stopping_rounds=200,
                   verbose_eval=100,
                   valid_sets=[train,train_test]
                   )

```

Training until validation scores don't improve for 200 rounds.

```

[100] training's multi_error: 0.0732298 valid_1's multi_error: 0.150007
[200] training's multi_error: 0.0513315 valid_1's multi_error: 0.0836788
[300] training's multi_error: 0.0446361 valid_1's multi_error: 0.0819834
[400] training's multi_error: 0.0376761 valid_1's multi_error: 0.0810803
[500] training's multi_error: 0.0339035 valid_1's multi_error: 0.0818256
Early stopping, best iteration is:
[392] training's multi_error: 0.0380341 valid_1's multi_error: 0.08098


```

In []:

```
lgb.plot_importance(model)
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f5bale27470>



In []:

```
lgb.create_tree_digraph(model)
```


Out[]:

In []:

```
lgb.plot_tree(model)
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f5b9dc9d128>



In []:

```
from google.colab import drive
drive.mount('/content/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly

Enter your authorization code:

.....

Mounted at /content/drive

In []:

```
import os
os.chdir('/content/drive/My Drive/ML case study/models')
```

In []:

```
import joblib
joblib.dump(model, 'Model_building_LGB(1).pkl')
```

Out[]:

```
['Model_building_LGB(1).pkl']
```

In []:

```
import joblib
model = joblib.load('Model_building_LGB(1).pkl')
```

In []:

```
os.chdir('/content')
```

In []:

```
!wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/84.0.4147.105 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-GB,en-US;q=0.9,en;q=0.8" --header="Referer: https://www.kaggle.com/" "https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/11835/224935/compressed/test.csv.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1597062299&Signature=sSEFif%2Bviyz4WtNlQ0YoOIKIdFcDpMHNbPp6WC0PyglcoixHb3mbDv90BDKUgfaGHRdlNL5jZdlYFSKTabezCGvWs2GLuvwjD5b2DyS6lOXgv1zJDoEYiZSaw2P3oCCRjSJKQKv5GZeHEXA4go%2FpMg4jzDFh%2BCnTxCPdFeYntOBu%2F8NgEe7ntcLUDvbtzn%2BwIrBuudlRDDVybuv2fNTuzoQLgDvM2U%2B6NzaYZVzkEbxvGJRCieT7b0RxegoC3LRpl%2Bk0WPXo6Qy4qq7quVNix%2FjvkjPhXz2txIj27sGcn%2Bg%2FGP8OKCvWitUAPnry7ngQLMxwdhyMASE4iXIkRl4dGA%3D%3D&response-content-disposition=attachment%3B+filename%3Dtest.csv.zip" -c -O 'test.csv.zip'
```

```
--2020-08-08 11:57:06-- https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/11835/224935/compressed/test.csv.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1597062299&Signature=sSEFif%2Bviyz4WtNlQ0YoOIKIdFcDpMHNbPp6WC0PyglcoixHb3mbDv90BDKUgfaGHRdlNL5jZdlYFSKTabezCGvWs2GLuvwjD5b2DyS6lOXgv1zJDoEYiZSaw2P3oCCRjSJKQKv5GZeHEXA4go%2FpMg4jzDFh%2BCnTxCPdFeYntOBu%2F8NgEe7ntcLUDvbtzn%2BwIrBuudlRDDVybuv2fNTuzoQLgDvM2U%2B6NzaYZVzkEbxvGJRCieT7b0RxegoC3LRpl%2Bk0WPXo6Qy4qq7quVNix%2FjvkjPhXz2txIj27sGcn%2Bg%2FGP8OKCvWitUAPnry7ngQLMxwdhyMASE4iXIkRl4dGA%3D%3D&response-content-disposition=attachment%3B+filename%3Dtest.csv.zip
```

Resolving storage.googleapis.com (storage.googleapis.com)... 173.194.214.128, 173.194.215.128, 173.194.216.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|173.194.214.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1791131386 (1.7G) [application/zip]
Saving to: 'test.csv.zip'

test.csv.zip 100%[=====>] 1.67G 135MB/s in 14s

2020-08-08 11:57:20 (124 MB/s) - 'test.csv.zip' saved [1791131386/1791131386]

In []:

```
unzip test.csv.zip
```

Archive: test.csv.zip
 inflating: test.csv

In []:

```
test = featureModify(False, None)  
print("Done test read")
```

Memory usage of dataframe is 3974.83 MB
Memory usage after optimization is: 1079.37 MB
Decreased by 72.8%
Done test read

In []:

```
test.head()
```

Out[]:

	id	crew	experiment	time	seat	eeg_fp1	eeg_f7	eeg_f8	eeg_t4	eeg_t6	eeg_t5	eeg_t3	eeg
0	0	1	LOFT	0.000000	0	17.906250	6.128906	0.994629	28.203125	47.687500	187.125000	33.187500	-4.22
1	1	1	LOFT	0.000000	1	45.875000	94.750000	23.296875	1.391602	2.060547	-5.144531	6.394531	33.40
2	2	1	LOFT	0.003906	0	33.125000	28.359375	-7.238281	-7.691406	25.828125	107.250000	12.843750	1.21
3	3	1	LOFT	0.003906	1	43.281250	95.875000	18.703125	-1.432617	-4.234375	-8.023438	7.425781	27.34
4	4	1	LOFT	0.007812	0	7.929688	3.460938	10.859375	26.359375	25.890625	37.000000	50.343750	11.67

In []:

```
df_sub = pd.DataFrame()  
df_sub['id'] = test['id']  
test = test.drop('id',axis=1)
```

In []:

```
test.head()
```

Out[]:

	crew	experiment	time	seat	eeg_fp1	eeg_f7	eeg_f8	eeg_t4	eeg_t6	eeg_t5	eeg_t3	eeg_fp2
0	1	LOFT	0.000000	0	17.906250	6.128906	0.994629	28.203125	47.687500	187.125000	33.187500	-4.222656
1	1	LOFT	0.000000	1	45.875000	94.750000	23.296875	1.391602	2.060547	-5.144531	6.394531	33.406250

2	crew	experiment	LOFT	time	seat	33.925000	28.553375	-7.253289	-7.659406	eeg t6	eeg t5	12.847438	12.847438
3	1	LOFT	0.003906	1	43.281250	95.875000	18.703125	-1.432617	-4.234375	-8.023438	7.425781	27.343750	
4	1	LOFT	0.007812	0	7.929688	3.460938	10.859375	26.359375	25.890625	37.000000	50.343750	11.679688	

In []:

```
import joblib
model = joblib.load('/content/drive/My Drive/ML case study/models/Model_building_LGB(1).pkl')
```

In []:

```
y_pred = model.predict(test,num_iteration=model.best_iteration)
```

In []:

```
y_pred
```

Out[]:

```
array([[0.97357333, 0.00215279, 0.02030727, 0.00396661],
       [0.96985874, 0.00287824, 0.02312566, 0.00413735],
       [0.97360533, 0.00219119, 0.02011564, 0.00408784],
       ...,
       [0.88274026, 0.00318973, 0.10606072, 0.00800928],
       [0.94154287, 0.00231302, 0.04893579, 0.00720831],
       [0.88331155, 0.00313693, 0.10554946, 0.00800207]])
```

In []:

```
import os
os.chdir('/content/drive/My Drive/ML case study/results')
```

In []:

```
import joblib
joblib.dump(y_pred, 'Model_building5_LGB_output(1).pkl')
```

Out[]:

```
['Model_building5_LGB_output(1).pkl']
```

In []:

```
import joblib
y_pred = joblib.load('/content/drive/My Drive/ML case study/results/Model_building5_LGB_output(1).pkl')
```

In []:

```
y_pred
```

Out[]:

```
array([[0.97357333, 0.00215279, 0.02030727, 0.00396661],
       [0.96985874, 0.00287824, 0.02312566, 0.00413735],
       [0.97360533, 0.00219119, 0.02011564, 0.00408784],
       ...,
       [0.88274026, 0.00318973, 0.10606072, 0.00800928],
       [0.94154287, 0.00231302, 0.04893579, 0.00720831],
       [0.88331155, 0.00313693, 0.10554946, 0.00800207]])
```

In []:

```
import os
os.chdir('/content/drive/My Drive/ML case study/results/')
```

In []:

```
df_sub = pd.DataFrame(np.concatenate((np.arange(len(test))[:, np.newaxis], y_pred), axis=1), columns=['id', 'A', 'B', 'C', 'D'])
df_sub['id'] = df_sub['id'].astype(int)
print(df_sub)
df_sub.to_csv("Model_building5_LGBop.csv", index=False)
```

	id	A	B	C	D
0	0	0.973573	0.002153	0.020307	0.003967
1	1	0.969859	0.002878	0.023126	0.004137
2	2	0.973605	0.002191	0.020116	0.004088
3	3	0.969847	0.002912	0.023105	0.004136
4	4	0.974038	0.002179	0.019839	0.003944
...
17965138	17965138	0.882729	0.003132	0.106111	0.008028
17965139	17965139	0.941786	0.002305	0.048648	0.007261
17965140	17965140	0.882740	0.003190	0.106061	0.008009
17965141	17965141	0.941543	0.002313	0.048936	0.007208
17965142	17965142	0.883312	0.003137	0.105549	0.008002

[17965143 rows x 5 columns]

Random Forest

In []:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
param_dist = {"n_estimators": [1, 5, 10],
              "max_depth" : [200],
              "criterion" : ['entropy'],
              "random_state" : [40],
              "max_features" : ['auto'],
              }
rfc = RandomForestClassifier()
rsc = RandomizedSearchCV(rfc, param_dist, cv=2, verbose=500)
rsc.fit(X_train, y_train)
print("Tuned Decision Tree Parameters: {}".format(rsc.best_params_))
print("Best score is {}".format(rsc.best_score_))
```

In []:

```
from sklearn.ensemble import RandomForestClassifier
from tqdm.auto import tqdm
n_estimators = [10, 25, 50]
for i in tqdm(n_estimators):
    print('n_estimators are:', i)
    clf = RandomForestClassifier(n_estimators=i, max_depth=200, criterion='entropy')
    clf.fit(X_train, y_train)
    y_pred = clf.predict_proba(X_test)
    print('Log_loss is:', (log_loss(y_test, y_pred)))
```

```
n_estimators are: 10
Log_loss is: 0.043098855808167656
n_estimators are: 25
Log_loss is: 0.033084208336728135
n_estimators are: 50
Log_loss is: 0.03091853122729293
```

In []: